

Lecture

CMPU4021

Distributed Systems

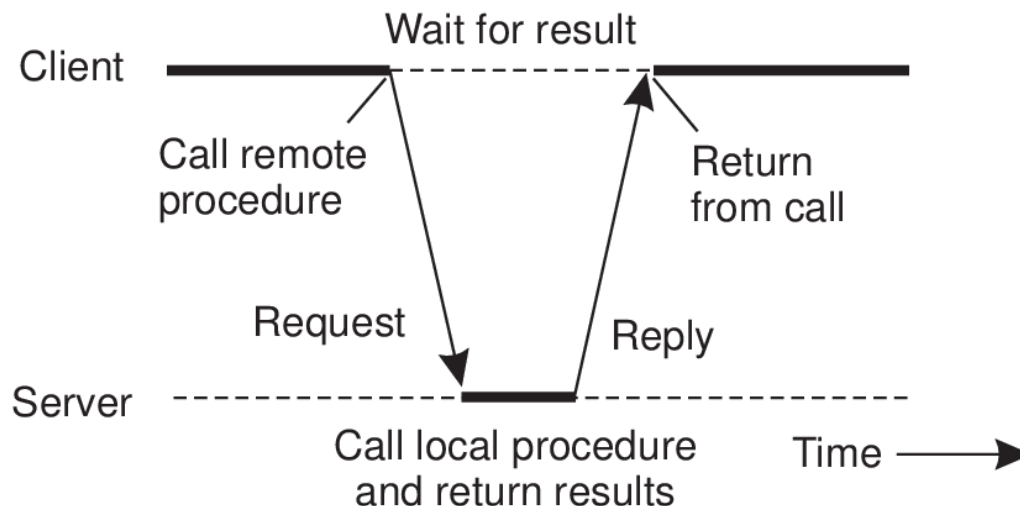
Remote Invocation

Remote Invocation

- Covers a range of techniques
- Based on a two-way exchange between communicating entities in a distributed system
 - resulting in the calling of a remote operation, procedure or method.
- Remote Procedure Calls (RPC)
- Remote method invocation (RMI)

Basic RPC operation

- Supports client-server computing with servers offering a set of operations through a service interface and clients calling these operations directly as if they were available locally.
- RPC systems offer (at a minimum)
 - access transparency
 - the calling procedure should not be aware that the called procedure is executing on a different machine or vice versa.
 - location transparency

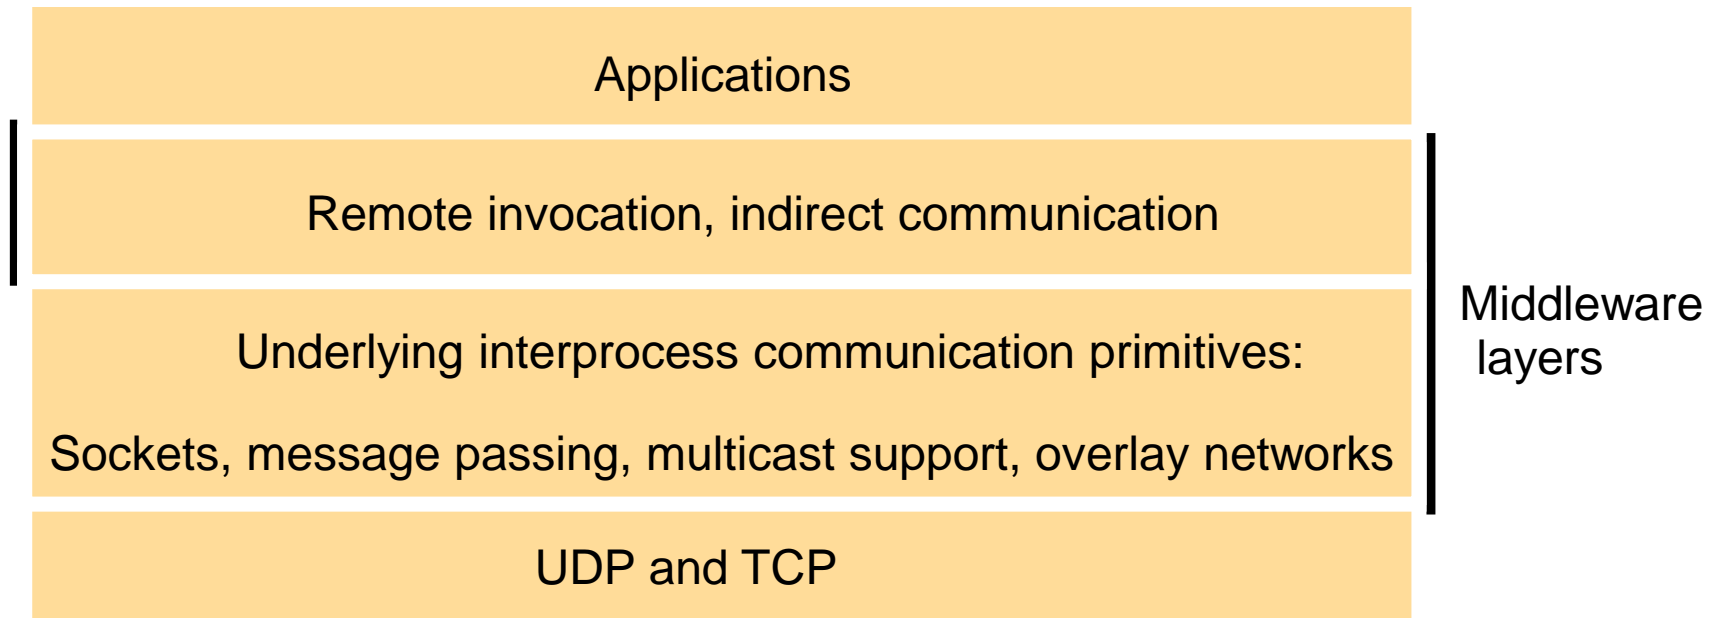


Remote method invocation (RMI)

- Strongly resembles remote procedure calls but in a world of distributed objects.
- A calling object can invoke a method in a remote object.
- Communication among distributed objects via RMI
 - Recipients of remote invocations are remote objects, which implement remote interfaces for communication
- Reliability
 - Either one or both the invoker and invoked can fail, and status of communication is supported by the interface. e.g.,
 - notification on failures,
 - reply generation,
 - parameter processing – marshalling/unmarshalling
- Local invocations target local objects, and remote invocations target remote objects

RMI in Middleware Layers

- A suite of API software that uses underlying processes and communication (message passing) protocols to provide its abstract protocol – simple RMI request-reply protocol



Interfaces

- Interfaces hide the details of modules providing the service; and access to module variables is only indirectly via 'getter' and 'setter' methods / mechanisms associated with the interfaces
 - e.g., call by value/reference for local calls through pointers vs. input, output, and input paradigms in RMI through message-data and objects
- *Service interfaces*
 - client-server model, specification of the procedures offered by a server
 - defining the types of input and output arguments
- *Remote interfaces*
 - distributed object model, specifies the methods of an object that are available for invocation by objects in other processes
 - defining the types of the input and output arguments of each of them.

Interface definition languages (IDLs)

- IDL is a language that is used to define the interface between a client and server process in a distributed system.
- Each interface definition language also has a set of associated IDL compilers
 - one per supported target language.
- An IDL compiler compiles the interface specifications, listed in an IDL input file, into source code (e.g., C/C++, Java) that implements the low-level communication details required to support the defined interfaces.
- Provides a notation for defining interfaces in which each of the parameters of a method may be described as for *input* or output in addition to having its type specified.

The distributed object model

- RMI
 - invocations between objects in different processes (either on same or different computers) is *remote*. Invocations within the same process are *local*
- Each process contains objects, some of which can receive remote invocations, other only local invocations
- Those that can receive remote invocations are called remote objects
- Objects need to know the remote object reference of an object in another process in order to invoke its methods. How do they get it?
 - The remote interface specifies which methods can be invoked remotely

Distributed Garbage Collection

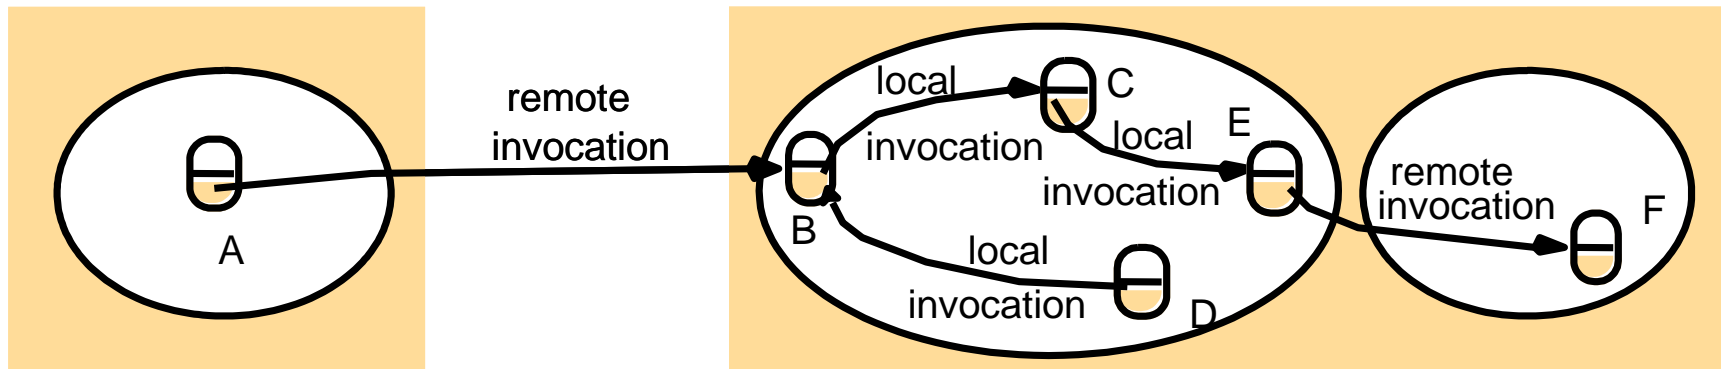
- Where any process includes remote objects, then it is equipped with both
 - Local garbage collector
 - Distributed garbage collector
- For any remote object `O`
 - `O.holders` is a list of all the processes that have a remote reference to that object i.e. got a stub for it
- When a client `C` receives a remote reference for `O` it makes an `addRef` call to `O`'s garbage collector
 - resulting in its being added to `O.holders`
- When `C`'s local garbage collector attempts to delete the stub object for `O`, it calls `removeRef` on `O`'s garbage collector, resulting
 - it its being removed from `O.holders`
- When `O.holders` is empty, `O` can be deleted

Distributed Garbage Collection

- Possible difficulties
 - `removeRef` and `addRef` sent at same time from different processes
- Possibility that O would be deleted, even though a client *thinks* it has a reference
- Incorporate a delay / temporary reference to solve
 - `addRef` goes missing
- Client must detect, exception returned
 - `removeRef` message goes missing / not sent
- Time based leases are allocated for objects

The distributed object model: Remote and local method invocations

- Objects receiving remote invocations (service objects) are remote objects, e.g., B and F
- Object references are required for invocation, e.g., C must have E's reference or B must have A's reference
- B and F must have remote interfaces (of their accessible methods)

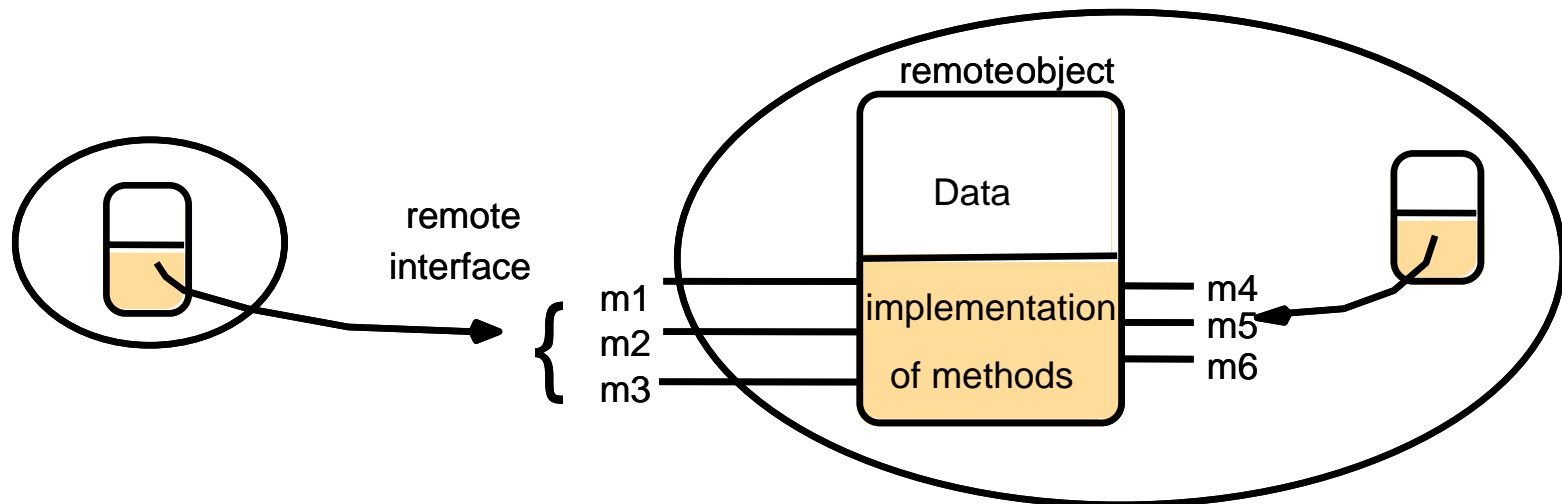


The distributed object model

- Remote object references
 - An unique identifier of a remote object, used throughout a distributed system
 - The remote object reference (including the 'interface' list of methods) can be passed as arguments or results in rmi.

The distributed object model: Remote interfaces

- Remote objects have a class that implement remote methods (as public).
- Local objects can access methods in an interface plus methods implemented by remote objects
 - Remote interfaces can't be constructed – no constructors



Invocation Semantics

- Unreliable network
 - For all request –reply protocols, messages may get lost
- Solutions for lost / retransmitted messages
 - Retry request
 - Filter Duplicates
 - Retransmit results

Invocation Semantics

- In local OO system
 - all methods are invoked exactly once per request –guaranteed – unless whole process fails
- In distributed object system, we need to know what has happened if we do not hear result from remote object i.e. did the request go missing, did the response go missing
- 3 different types of guarantee (invocation semantics) may be provided
 - could be implemented in a middleware platform intended to support remote method invocations:
 - *Maybe*
 - *At-Least-Once*
 - *At-Most-Once*

Maybe Invocation Semantics

- If the invoker cannot tell whether a remote method has been invoked or not
- Very inexpensive, but only useful if the system can tolerate occasional failed invocations

At-Least-Once

Invocation Semantics

- If the invoker receives a result, then it is guaranteed that the method was invoked at least once
- Achieved by resending requests to mask omission failure
- Only useful if the operations are idempotent ($x = 10$, rather than $x = x + 10$)
- Inexpensive on server

At-Most-Once

Invocation Semantics

- If the invoker receives a result, then it is guaranteed that the method was invoked only once
- If no result is received, then the method was executed either never or once
- Achieved by resending requests, and storing and resending responses
- More expensive on a server / remote object, which must maintain results and recognise duplicate messages

Invocation semantics: failure model

- Maybe, At-least-once and At-most-once
 - can suffer from crash failures when the server containing the remote object fails.
- *Maybe*
 - if no reply, the client does not know if method was executed or not
 - omission failures - if the invocation or result message is lost
- *At-least-once*
 - the client gets a result (and the method was executed at least once) or an exception (no result)
 - arbitrary failures. If the invocation message is retransmitted, the remote object may execute the method more than once, possibly causing wrong values to be stored or returned.
 - if *idempotent* operations are used, arbitrary failures will not occur
- *At-most-once*
 - the client gets a result (and the method was executed exactly once) or an exception (instead of a result, in which case, the method was executed once or not at all)

Invocation semantics: failure model

<i>Fault tolerance measures</i>			<i>Invocation semantics</i>
<i>Retransmit request message</i>	<i>Duplicate filtering</i>	<i>Re-execute procedure or retransmit reply</i>	
No	Not applicable	Not applicable	<i>Maybe</i>
Yes	No	Re-execute procedure	<i>At-least-once</i>
Yes	Yes	Retransmit reply	<i>At-most-once</i>

Design Issues of RMI

- Local invocations have
 - at-most-once, or
 - exactly-once semantics
- Distributed RMI, the alternative invocation semantics are:
 - *Retry request message* – retransmit until reply is received or on server failure – at-least-once semantics;
 - *Duplicate message filtering* – discard duplicates at server (using seq #s or ReqID);
 - Buffer result messages at server for *retransmission* – avoids redo of requests (even for idempotent ops) – at-most-once semantics.
 - Idempotent operation – the one which can be performed repeatedly with the same effect as if it had been performed exactly once.

Transparency

- Remote invocations should be made transparent
 - the syntax of a remote invocation is the same as that of a local invocation, but
 - the difference between local and remote objects should be expressed in their interfaces.
- E.g. Java RMI:
 - Remote objects implement `Remote` interface and throw *RemoteExceptions*
- Remote object should be able to keep its state consistent in the presence of concurrent accesses from multiple clients.

Representation of a remote object reference

- a remote object reference must be unique in the distributed system and over time. It should not be reused after the object is deleted.

Why not?

- the first two fields locate the object unless migration or re-activation in a new process can happen
- the fourth field identifies the object within the process
- its interface tells the receiver what methods it has (e.g. class *Method*)
- a remote object reference is created by a remote reference module when a reference is passed as argument or result to another process
 - it will be stored in the corresponding proxy
 - it will be passed in request messages to identify the remote object whose method is to be invoked

32 bits	32 bits	32 bits	32 bits	
Internet address	port number	time	object number	interface of ₂₃ remote object

Distributed garbage collection

Ensures

- if a local or remote reference to an object is still held anywhere in a set of distributed objects
 - then the object itself will continue to exist,
- As soon as no object holds a reference to it
 - the object is collected and the memory it uses recovered.

Summary

RMI

- Each object has a (global) remote object reference and a remote interface that specifies which of its operations can be invoked remotely.
- Local method invocations provide exactly-once semantics
 - the best RMI can guarantee is at-most-once
- Middleware components proxies, skeletons and dispatchers hide details of marshalling, message passing and object location from programmers.

References

- Chapter 5: Coulouris, Dollimore and Kindberg, Distributed Systems: Concepts and Design, 5/E
- Chapter 4: Maarten van Steen, Andrew S. Tanenbaum Distributed Systems, 3rd edition (2017)